

```
In [202]: ## Imports for the dataset and building the neural networks.
import nltk
from nltk.corpus import reuters
import keras
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import sequence
from keras.preprocessing.image import ImageDataGenerator
from keras.regularizers import l2
from keras import backend as K
from keras.models import Sequential, load_model
from keras.layers.core import Dense, Dropout, Activation
from keras.utils import np_utils
from keras.utils import to_categorical
from keras.utils import plot_model
from keras import models
from keras import optimizers
from keras import layers
from keras.datasets import reuters
from keras.models import Sequential
from keras.layers import SimpleRNN
from keras.layers import LSTM
from keras.layers import MaxPooling2D
from keras.layers import Dense
from keras.layers import Embedding
from keras.layers import Dropout
from keras.layers import Dense, Conv2D, BatchNormalization, Activation
from keras.layers import AveragePooling2D, Input, Flatten
from keras.preprocessing import image
from keras.preprocessing.image import ImageDataGenerator
from keras.callbacks import ModelCheckpoint
from keras.callbacks import EarlyStopping
from keras.callbacks import ReduceLROnPlateau
from keras.callbacks import LearningRateScheduler
from keras.optimizers import Adam
from keras import regularizers
from keras.applications import VGG16
import pickle
import tensorflow as tf
from sklearn import metrics
from sklearn.model_selection import train_test_split
import numpy as np
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import os
from urllib.request import urlretrieve
from os.path import isfile, isdir
from tqdm import tqdm
import tarfile
import json
from time import time
import chakin
```

Hyperparameters

- `max_features` and `maxlen` : Cuts off texts after this many words amongst most common words
- `max_words` : Number of words to consider as features
- `epochs` : number of iterations until the network stops learning or start overfitting
- `batch_size` : highest number that your machine has memory for. Most people set them to common sizes of memory:
- `learning_rate` : number how fast the model learns

```
In [203]: maxlen=1000 # Cuts off texts after this many words amongst most com
mon words
max_features=1000
max_words=10000 # Number of words to consider as features
#max_features=10000
#max_words=10000
#maxlen=10000
batch_size=128
epochs = 10
learning_rate = 0.001
```

```

In [204]: ##https://towardsdatascience.com/text-classification-in-keras-part-1-a-simple-reuters-news-classifier-9558d34d01d3
# save np.load
np_load_old = np.load

# modify the default parameters of np.load
np.load = lambda *a,**k: np_load_old(*a, allow_pickle=True, **k)

#Was working: (trainingDataRaw, trainingLabelsRaw), (testingDataRaw, testingLabelsRaw) = reuters.load_data(num_words=None, test_split=0.2)
(trainingDataRaw, trainingLabelsRaw), (testingDataRaw, testingLabelsRaw) = reuters.load_data(num_words=max_features, test_split=0.2)
# restore np.load for future normal usage
np.load = np_load_old
word_index = reuters.get_word_index(path="reuters_word_index.json")

print('# of Training Samples: {}'.format(len(trainingDataRaw)))
print('# of Test Samples: {}'.format(len(testingDataRaw)))

num_classes = max(trainingLabelsRaw) + 1
print('# of Classes: {}'.format(num_classes))
# of Training Samples: 8982
# of Test Samples: 2246
# of Classes: 46
index_to_word = {}
for key, value in word_index.items():
    index_to_word[value] = key
print(' '.join([index_to_word[x] for x in trainingDataRaw[0]]))
print(trainingLabelsRaw[0])

# of Training Samples: 8982
# of Test Samples: 2246
# of Classes: 46
the of of mln loss for plc said at only ended said of could 1 tra
ders now april 0 a after said from 1985 and from foreign 000 apri
l 0 prices its account year a but in this mln home an states earl
ier and rise and revs vs 000 its 16 vs 000 a but 3 of of several
and shareholders and dividend vs 000 its all 4 vs 000 1 mln agree
d of april 0 are 2 states will billion total and against 000 pct
dlrs
3

```

```
In [205]: from keras.preprocessing.text import Tokenizer

tokenizer = Tokenizer(num_words=max_features)
trainingData = tokenizer.sequences_to_matrix(trainingDataRaw, mode='binary')
testingData = tokenizer.sequences_to_matrix(testingDataRaw, mode='binary')

trainingLabels = keras.utils.to_categorical(trainingLabelsRaw, num_classes)
testingLabels = keras.utils.to_categorical(testingLabelsRaw, num_classes)
print(trainingData[0])
print(len(trainingData[0]))

print(trainingLabels[0])
print(len(trainingLabels[0]))
```

[illegible]

```

In [206]: CHAKIN_INDEX = 11
          NUMBER_OF_DIMENSIONS = 50
          SUBFOLDER_NAME = "glove.6B"

          DATA_FOLDER = "embeddings"
          ZIP_FILE = os.path.join(DATA_FOLDER, "{}.zip".format(SUBFOLDER_NAME))
          ZIP_FILE_ALT = "glove" + ZIP_FILE[5:] # sometimes it's lowercase only...
          UNZIP_FOLDER = os.path.join(DATA_FOLDER, SUBFOLDER_NAME)
          if SUBFOLDER_NAME[-1] == "d":
              GLOVE_FILENAME = os.path.join(
                  UNZIP_FOLDER, "{}.txt".format(SUBFOLDER_NAME))
          else:
              GLOVE_FILENAME = os.path.join(UNZIP_FOLDER, "{}.{}d.txt".format(
                  SUBFOLDER_NAME, NUMBER_OF_DIMENSIONS))

          if not os.path.exists(ZIP_FILE) and not os.path.exists(UNZIP_FOLDER):
              # GloVe by Stanford is licensed Apache 2.0:
              # https://github.com/stanfordnlp/GloVe/blob/master/LICENSE
              # http://nlp.stanford.edu/data/glove.twitter.27B.zip
              # Copyright 2014 The Board of Trustees of The Leland Stanford Junior University
              print("Downloading embeddings to {}".format(ZIP_FILE))
              chakin.download(number=CHAKIN_INDEX, save_dir='{}'.format(DATA_FOLDER))
          else:
              print("Embeddings already downloaded.")

          if not os.path.exists(UNZIP_FOLDER):
              import zipfile
              if not os.path.exists(ZIP_FILE) and os.path.exists(ZIP_FILE_ALT):
                  ZIP_FILE = ZIP_FILE_ALT
              with zipfile.ZipFile(ZIP_FILE, "r") as zip_ref:
                  print("Extracting embeddings to {}".format(UNZIP_FOLDER))
                  zip_ref.extractall(UNZIP_FOLDER)
          else:
              print("Embeddings already extracted.")

```

Embeddings already downloaded.

Embeddings already extracted.

```

In [207]: def load_embedding_from_disks(embeddings_filename, with_indexes=True):
    """
    Read a embeddings txt file. If `with_indexes=True`,
    we return a tuple of two dictionaries
    `(word_to_index_dict, index_to_embedding_array)`,
    otherwise we return only a direct
    `word_to_embedding_dict` dictionary mapping
    from a string to a numpy array.
    """
    if with_indexes:
        word_to_index_dict = dict()
        index_to_embedding_array = []

    else:
        word_to_embedding_dict = dict()

    with open(embeddings_filename, 'r', encoding='utf-8') as embeddings_file:
        for (i, line) in enumerate(embeddings_file):

            split = line.split(' ')

            word = split[0]

            representation = split[1:]
            representation = np.array(
                [float(val) for val in representation]
            )

            if with_indexes:
                word_to_index_dict[word] = i
                index_to_embedding_array.append(representation)
            else:
                word_to_embedding_dict[word] = representation

    # Empty representation for unknown words.
    _WORD_NOT_FOUND = [0.0] * len(representation)
    if with_indexes:
        _LAST_INDEX = i + 1
        word_to_index_dict = defaultdict(
            lambda: _LAST_INDEX, word_to_index_dict)
        index_to_embedding_array = np.array(
            index_to_embedding_array + [_WORD_NOT_FOUND])
        return word_to_index_dict, index_to_embedding_array
    else:
        word_to_embedding_dict = defaultdict(lambda: _WORD_NOT_FOUND)

    return word_to_embedding_dict

print('\nLoading embeddings from', GLOVE_FILENAME)
word_to_index, index_to_embedding = \
    load_embedding_from_disks(GLOVE_FILENAME, with_indexes=True)
print("Embedding loaded from disks.")

```

```
Loading embeddings from embeddings/glove.6B/glove.6B.50d.txt  
Embedding loaded from disks.
```



```

In [208]: # Compare training and dev
def plot_history(history):
    accuracy = history.history['acc']
    val_accuracy = history.history['val_acc']
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    epoch_number = range(1, len(accuracy) + 1)
    plt.style.use('ggplot') # Grammar of Graphics plots
    plt.figure(figsize=(20, 10))
    plt.subplot(1, 2, 1)
    plt.plot(epoch_number, accuracy, 'b', label='Training')
    plt.plot(epoch_number, val_accuracy, 'r', label='Dev')
    plt.title('Training and Dev Set Accuracy')
    plt.xlabel('Epoch Number')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.subplot(1, 2, 2)
    plt.plot(epoch_number, loss, 'b', label='Training')
    plt.plot(epoch_number, val_loss, 'r', label='Dev')
    plt.title('Training and Dev Set Loss')
    plt.xlabel('Epoch Number')
    plt.ylabel('Loss')
    plt.legend()
    plt.savefig('fig-training-process.pdf',
                papertype = 'letter', orientation = 'landscape')
    plt.show()
    plt.close()

# plot confusion matrix to external file
def plot_confusion(cm_data):
    plt.figure()
    selected_cmap = sns.cubehelix_palette(light=1, as_cmap=True)
    sns_plot = sns.heatmap(cm_data, annot=True, fmt="d", \
                            cmap = selected_cmap, linewidths = 0.5, cbar = False)
    sns_plot.set_yticklabels(sns_plot.get_yticklabels(), rotation =
0)
    plt.title('Confusion Matrix')
    plt.ylabel('True Digit')
    plt.xlabel('Predicted Digit')
    # plt.show() # use if plot to screen is desired
    plt.savefig('fig-confusion-matrix.pdf',
                papertype = 'letter', orientation = 'landscape')
    plt.close()

def train_neural_network(session, optimizer, keep_probability, feature_batch, label_batch):
    session.run(optimizer,
                  feed_dict={
                      x: feature_batch,
                      y: label_batch,
                      keep_prob: keep_probability
                  })

def print_stats(session, feature_batch, label_batch, cost, accurac

```

```

In [209]: def normalize(x):
            """
            argument
            - x: input image data in numpy array [32, 32, 3]
            return
            - normalized x
            """
            min_val = np.min(x)
            max_val = np.max(x)
            x = (x-min_val) / (max_val-min_val)
            return x

def one_hot_encode(x):
    """
    argument
    - x: a list of labels
    return
    - one hot encoding matrix (number of labels, number of
class)
    """
    encoded = np.zeros((len(x), 47))

    for idx, val in enumerate(x):
        encoded[idx][val] = 1

    return encoded

def evaluate_model(model, train_features, train_labels, test_features, test_labels):
    # evaluate fitted model on the full training set
    train_loss, train_acc = model.evaluate(train_features, train_labels, verbose = 3)
    print('\nFull training set accuracy:', \
          '{:6.4f}'.format(np.round(train_acc, decimals = 4)))
    # evaluate the fitted model on the hold-out test set
    test_loss, test_acc = model.evaluate(test_features, test_labels, verbose = 3)
    print('Hold-out test set accuracy:', \
          '{:6.4f}'.format(np.round(test_acc, decimals = 4)))

def load_label_names():
    return ['cocoa', 'grain', 'veg-oil', 'earn', 'acq', 'wheat', 'copper',
            'housing', 'money-supply',
            'coffee', 'sugar', 'trade', 'reserves', 'ship', 'cotton', 'carcass',
            'crude', 'nat-gas',
            'cpi', 'money-fx', 'interest', 'gnp', 'meal-feed', 'alum',
            'oilseed', 'gold', 'tin',
            'strategic-metal', 'livestock', 'retail', 'ipi', 'iron-steel',
            'rubber', 'heat', 'jobs',
            'lei', 'bop', 'zinc', 'orange', 'pet-chem', 'dlr', 'gas', 'silver',
            'wpi', 'hog', 'lead']

def classification_report(model, test_features, test_labels):
    # examine the predicted values within a precision/recall framework

```

Targets:

```
Out[209]: ['cocoa',  
            'grain',  
            'veg-oil',  
            'earn',  
            'acq',  
            'wheat',  
            'copper',  
            'housing',  
            'money-supply',  
            'coffee',  
            'sugar',  
            'trade',  
            'reserves',  
            'ship',  
            'cotton',  
            'carcass',  
            'crude',  
            'nat-gas',  
            'cpi',  
            'money-fx',  
            'interest',  
            'gnp',  
            'meal-feed',  
            'alum',  
            'oilseed',  
            'gold',  
            'tin',  
            'strategic-metal',  
            'livestock',  
            'retail',  
            'ipi',  
            'iron-steel',  
            'rubber',  
            'heat',  
            'jobs',  
            'lei',  
            'bop',  
            'zinc',  
            'orange',  
            'pet-chem',  
            'dlr',  
            'gas',  
            'silver',  
            'wpi',  
            'hog',  
            'lead']
```

Model with no embedding space

Word vectors obtained from one-hot encoding.

Using the in-built keras functions that strip special characters and take into account the N most important words. These are high-dimensional and sparse.

Quite high accuracy since this is a word association problem and the semantic relationships aren't as important.

The resulting space has no structure.

```
In [210]: noEmbeddingModel = Sequential()
noEmbeddingModel.add(Dense(512, input_shape=(max_features,)))
noEmbeddingModel.add(Activation('relu'))
noEmbeddingModel.add(Dropout(0.3))
noEmbeddingModel.add(Dense(num_classes))
noEmbeddingModel.add(Activation('softmax')) #Softmax is used for multi-class classification
noEmbeddingModel.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
noEmbeddingModel.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_60 (Dense)	(None, 512)	512512
activation_31 (Activation)	(None, 512)	0
dropout_25 (Dropout)	(None, 512)	0
dense_61 (Dense)	(None, 46)	23598
activation_32 (Activation)	(None, 46)	0
=====		
Total params: 536,110		
Trainable params: 536,110		
Non-trainable params: 0		
=====		

```
In [211]: testingData.shape
```

```
Out[211]: (2246, 1000)
```

```
In [212]: model_name = "NoEmbeddingModel.h5"
train_model(noEmbeddingModel, model_name, trainingData, trainingLabels, testingData, testingLabels)
```

Training...

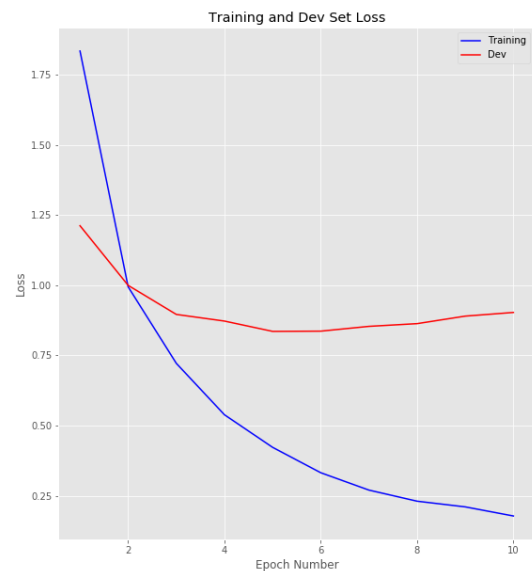
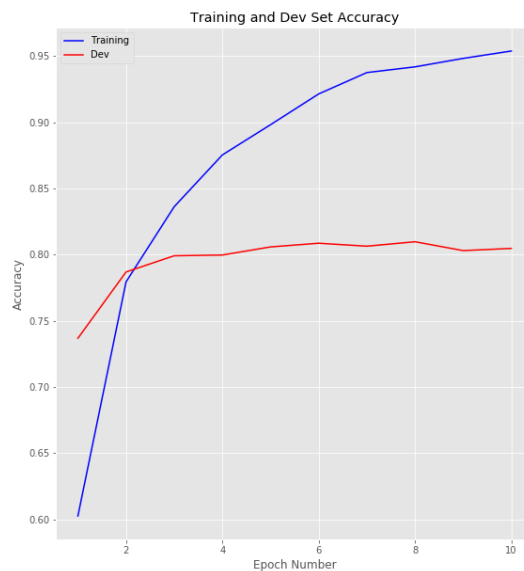
Epoch 00010: early stopping

Saved best trained model at NoEmbeddingModel.h5

Time of execution for training (seconds): 15.222

Full training set accuracy: 0.9334

Hold-out test set accuracy: 0.7890



Out[212]: <keras.callbacks.History at 0x7f2e62a8b8d0>

Model with no embedding space: Working with a small number of cases

Surprising high accuracy still. High training set accuracy due to overfitting.

```
In [213]: model_name="NoEmbeddingSmall.h5"
train_model(noEmbeddingModel, model_name, trainingData[-100:], trainingLabels[-100:], testingData, testingLabels)
```

Training...

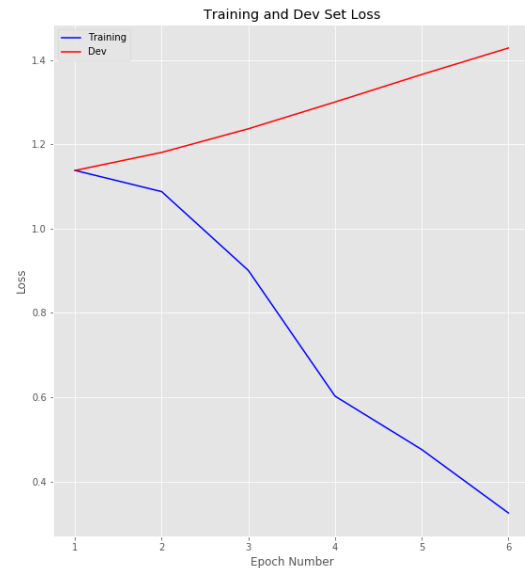
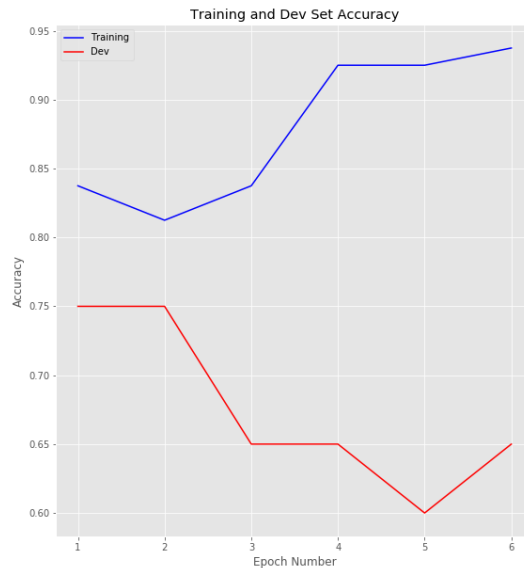
Epoch 00006: early stopping

Saved best trained model at NoEmbeddingSmall.h5

Time of execution for training (seconds): 0.069

Full training set accuracy: 0.9100

Hold-out test set accuracy: 0.7752



Out[213]: <keras.callbacks.History at 0x7f2dc881bbe0>

Model with Learned Embedding Space

Learned embeddings that are low-dimensional floating-point vectors that are learned from these data.

Some semantic relationships between the words are coded as vectors.

```
In [214]: learnedEmbeddingModel = Sequential()

#learnedEmbeddingModel.add(Embedding(max_features, 8, input_length=
maxlen))
learnedEmbeddingModel.add(Embedding(max_features, embedding_dim, in
put_length=maxlen))
learnedEmbeddingModel.add(Flatten())
learnedEmbeddingModel.add(Dense(num_classes))
learnedEmbeddingModel.add(Activation('relu'))
learnedEmbeddingModel.add(Dropout(0.3))
learnedEmbeddingModel.add(Dense(num_classes, activation='softmax'))
learnedEmbeddingModel.compile(optimizer='adam', loss='categorical_c
rossentropy', metrics=['acc'])
learnedEmbeddingModel.summary()
```

Layer (type)	Output Shape	Param #
embedding_25 (Embedding)	(None, 1000, 50)	50000
flatten_14 (Flatten)	(None, 50000)	0
dense_62 (Dense)	(None, 46)	2300046
activation_33 (Activation)	(None, 46)	0
dropout_26 (Dropout)	(None, 46)	0
dense_63 (Dense)	(None, 46)	2162
Total params: 2,352,208		
Trainable params: 2,352,208		
Non-trainable params: 0		

```
In [215]: model_name = "LearnedEmbedding.h5"
train_model(learnedEmbeddingModel, model_name, trainingData, trainingLabels, testingData, testingLabels)
```

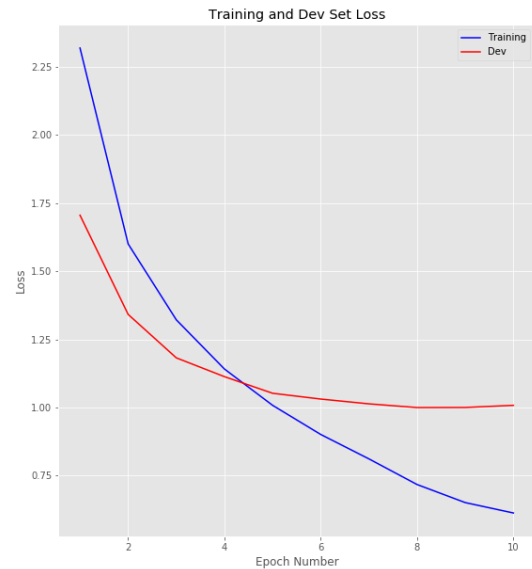
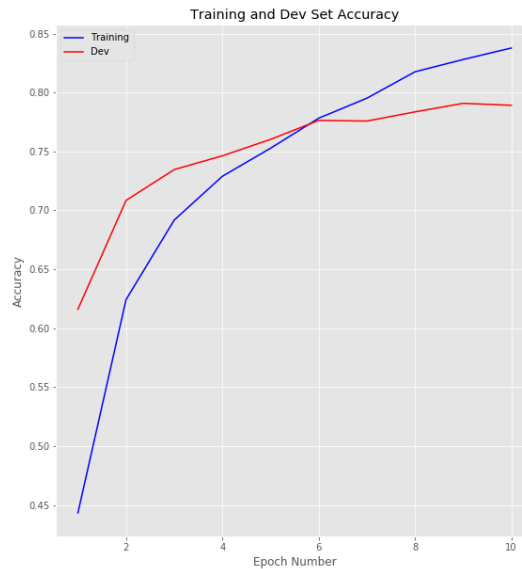
Training...

Saved best trained model at LearnedEmbedding.h5

Time of execution for training (seconds): 22.840

Full training set accuracy: 0.8898

Hold-out test set accuracy: 0.7796



Out[215]: <keras.callbacks.History at 0x7f2e180af160>

Model with Pre-Trained Embedding Layer

The pretrained embedding layer captures the generic aspects of the language

```
In [216]: #https://medium.com/@sabber/classifying-yelp-review-comments-using-
cnn-lstm-and-pre-trained-glove-word-embeddings-part-3-53fcea9a17fa
embeddings_index = dict()
f = open(GLOVE_FILENAME)
for line in f:
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs
f.close()
```



```
In [217]: embedding_dim = 50

embedding_matrix = np.zeros((max_words, embedding_dim))
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word)
    if i < max_words:
        if embedding_vector is not None:
            # Words not found in embedding index will be all-zeros.
            embedding_matrix[i] = embedding_vector
```

```
In [218]: pretrainedModel = Sequential()
pretrainedModel.add(Embedding(max_words, embedding_dim, input_length=maxlen, weights=[embedding_matrix], trainable=False))
pretrainedModel.add(Flatten())
pretrainedModel.add(Dense(32, activation='relu'))
pretrainedModel.add(Dropout(0.3))
pretrainedModel.add(Dense(46, activation='softmax'))
pretrainedModel.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc'])
pretrainedModel.summary()
```

Layer (type)	Output Shape	Param #
=====		
embedding_26 (Embedding)	(None, 1000, 50)	500000
flatten_15 (Flatten)	(None, 50000)	0
dense_64 (Dense)	(None, 32)	1600032
dropout_27 (Dropout)	(None, 32)	0
dense_65 (Dense)	(None, 46)	1518
=====		
Total params: 2,101,550		
Trainable params: 1,601,550		
Non-trainable params: 500,000		
=====		

```
In [219]: model_name="pretrained.h5"
          train_model(pretrainedModel, model_name, trainingData, trainingLabels, testingData, testingLabels)
```

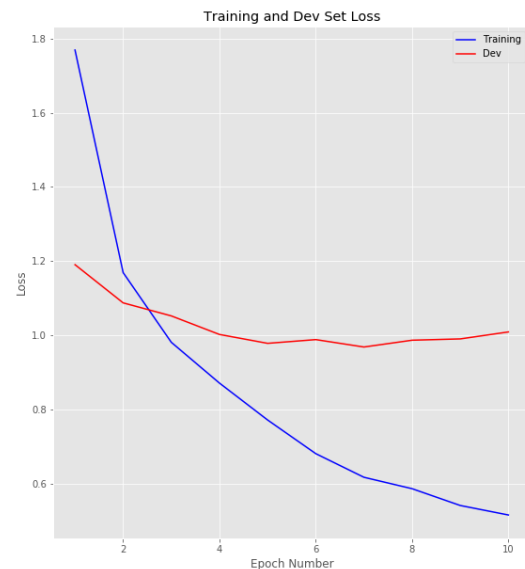
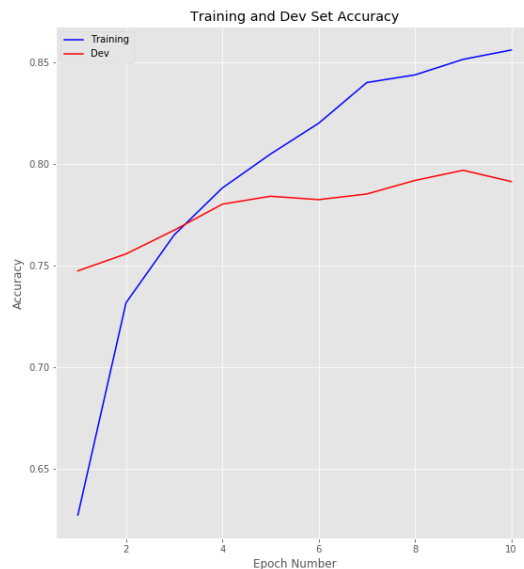
Training...

Saved best trained model at pretrained.h5

Time of execution for training (seconds): 20.107

Full training set accuracy: 0.9050

Hold-out test set accuracy: 0.7738



Out[219]: <keras.callbacks.History at 0x7f2e6ce770f0>

Pre-Trained Models: Working with a small number of cases

Pre-Trained embeddings are useful when there aren't enough data to learn the features

This shows similar accuracy as above which makes sense since the embeddings help with the learning.

The embeddings are computed using word-occurrence statistics (what words co-occur in sentences or documents)

<https://nlp.stanford.edu/projects/glove> (<https://nlp.stanford.edu/projects/glove>). Developed by Stanford researchers in 2014 based on factorizing a matrix of

word co-occurrences statistics.

```
In [220]: model_name="pretrainedSmall.h5"
          train_model(pretrainedModel, model_name, trainingData[-100:], train
                    ingLabels[-100:], testingData, testingLabels)
```

Training...

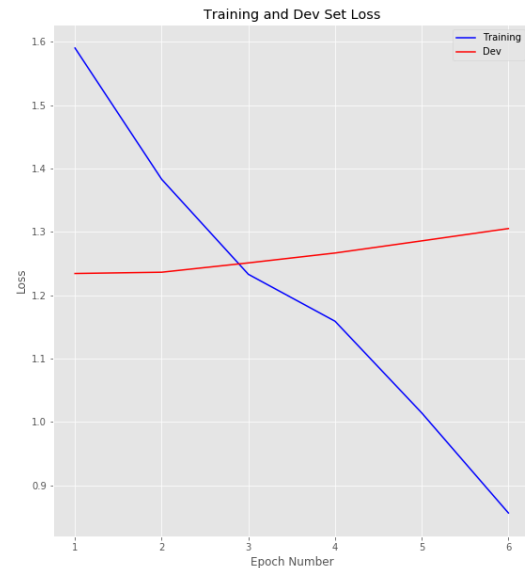
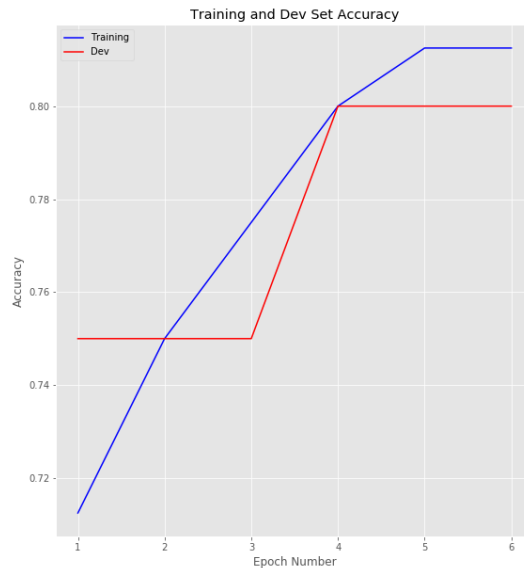
Epoch 00006: early stopping

Saved best trained model at pretrainedSmall.h5

Time of execution for training (seconds): 0.152

Full training set accuracy: 0.9000

Hold-out test set accuracy: 0.7663



Out[220]: <keras.callbacks.History at 0x7f2e185fbcc0>

SimpleRNNs

Processes sequences by iterating through the sequence elements and maintaining a state containing

information relative to what it has seen so far.

Takes inputs of (batch_size, timesteps, input_features)

SimpleRNNs aren't good at processing long sequences, such as text. Due to the vanishing gradient problem.

Refer to Hochreiter and Schmidhuber (1997).

SimpleRNNs aren't very good at processing long sequences, like text. Other types of recurrent layers perform much better. Let's take a look at some more advanced layers.

```
In [221]: embedding_matrix = np.zeros((max_words, embedding_dim))
          for word, i in word_index.items():
              embedding_vector = embeddings_index.get(word)
              if i < max_words:
                  if embedding_vector is not None:
                      # Words not found in embedding index will be all-zeros.
                      embedding_matrix[i] = embedding_vector

          trainingDataSeq = sequence.pad_sequences(trainingDataRaw, maxlen=maxlen)
          testingDataSeq = sequence.pad_sequences(testingDataRaw, maxlen=maxlen)
          print('input_train shape:', trainingDataSeq.shape)
          print('input_test shape:', testingDataSeq.shape)

input_train shape: (8982, 1000)
input_test shape: (2246, 1000)
```

```
In [222]: simpleRNNModel = Sequential()
simpleRNNModel.add(Embedding(max_words, embedding_dim, input_length
=maxlen, weights=[embedding_matrix], trainable=False))
simpleRNNModel.add(SimpleRNN(32, return_sequences=True))
simpleRNNModel.add(SimpleRNN(32, return_sequences=True))
simpleRNNModel.add(SimpleRNN(32, return_sequences=True))
simpleRNNModel.add(SimpleRNN(32))
simpleRNNModel.add(Dense(num_classes, activation='softmax'))
simpleRNNModel.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['acc'])

simpleRNNModel.summary()
```

Layer (type)	Output Shape	Param #
=====		
embedding_27 (Embedding)	(None, 1000, 50)	500000
<hr/>		
simple_rnn_33 (SimpleRNN)	(None, 1000, 32)	2656
<hr/>		
simple_rnn_34 (SimpleRNN)	(None, 1000, 32)	2080
<hr/>		
simple_rnn_35 (SimpleRNN)	(None, 1000, 32)	2080
<hr/>		
simple_rnn_36 (SimpleRNN)	(None, 32)	2080
<hr/>		
dense_66 (Dense)	(None, 46)	1518
=====		
Total params: 510,414		
Trainable params: 10,414		
Non-trainable params: 500,000		
<hr/>		

```
In [223]: model_name="simpleRNN.h5"
          train_model(simpleRNNModel, model_name, trainingDataSeq, trainingLabels,
                    testingDataSeq, testingLabels)
```

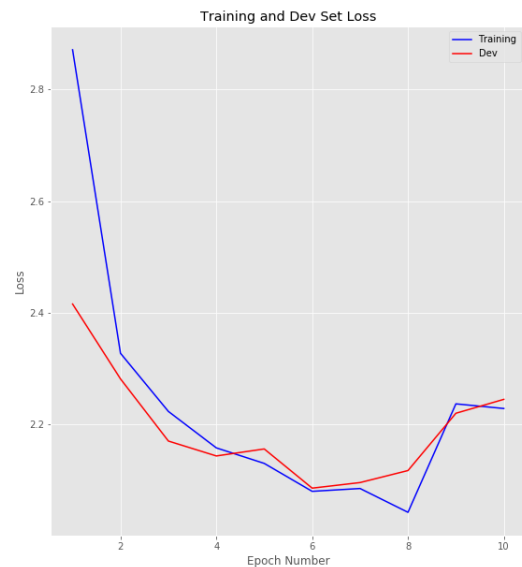
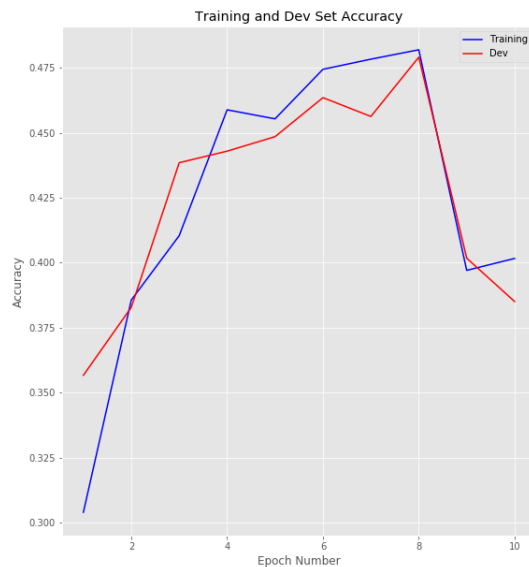
Training...

Saved best trained model at simpleRNN.h5

Time of execution for training (seconds): 1402.606

Full training set accuracy: 0.4006

Hold-out test set accuracy: 0.3931



Out[223]: <keras.callbacks.History at 0x7f2e67a4f128>

Long Short-Term Memory RNNs

LSTMs save information from the sequence to be ported later at a different point. So, it prevents older signals from vanishing.

Accuracy is much better than the SimpleRNN. I only trained using 3 epochs since it took so long to train.

Not performing better because the long-term structure of the text isn't as important with this problem because we're more concerned with word counts and mapping to the categories in this problem. The fully-connected approach does this fine.

The strength of LSTMs are more apparent with other problems like question-answering and machine translation.

```
In [224]: LSTMModel = Sequential()
LSTMModel.add(Embedding(max_words, embedding_dim, input_length=maxl
en, weights=[embedding_matrix], trainable=False))
LSTMModel.add(LSTM(32))
LSTMModel.add(Dense(num_classes, activation='softmax'))
LSTMModel.compile(optimizer='adam', loss='categorical_crossentropy
', metrics=['acc'])
LSTMModel.summary()
```

Layer (type)	Output Shape	Param #
embedding_28 (Embedding)	(None, 1000, 50)	500000
lstm_4 (LSTM)	(None, 32)	10624
dense_67 (Dense)	(None, 46)	1518

=====
 Total params: 512,142
 Trainable params: 12,142
 Non-trainable params: 500,000
 =====

```
In [225]: model_name="LSTM.h5"
#train_model(LSTMModel, model_name, trainingDataSeq, trainingLabel
s, testingDataSeq, testingLabels)
```

Much better accuracy since LSTM suffers less from the vanishing-gradient problem. Slightly better than the full-connected approach.

Not performing better as hyperparameters not tuned and I didn't include regularization.

Problem is well-solved looking at the frequency of words in the text when matching to the categories.

... this is what the fully-connected solution does.

RNNs and, in particular, LSTMs are better at machine translation and question-answering.

```
In [226]: begin_time = time()

## Compile and run the dense model
checkpointer = ModelCheckpoint(filepath=model_name,
                               monitor = 'val_acc',
                               verbose=2,
                               save_best_only=True)

early_stopping = EarlyStopping(monitor='val_loss',
                               min_delta=0,
                               patience=2,
                               verbose=2, mode='auto')

print('Training...')
## Training model
history = LSTMModel.fit(trainingDataSeq, trainingLabels,
                        batch_size=batch_size, epochs=epochs,
                        callbacks=[checkpointer, early_stopping],
                        verbose=2,
                        validation_split=0.2)

print('Saved best trained model at %s ' % model_name)

execution_time = time() - begin_time
print('\nTime of execution for training (seconds):', \
      '{:10.3f}'.format(np.round(execution_time, decimals = 3)))
evaluate_model(LSTMModel, trainingDataSeq, trainingLabels, testingDataSeq, testingLabels)
plot_history(history)
```


Training...

Train on 7185 samples, validate on 1797 samples

Epoch 1/10

- 100s - loss: 2.9474 - acc: 0.3488 - val_loss: 2.4200 - val_acc: 0.3450

Epoch 00001: val_acc improved from -inf to 0.34502, saving model to LSTM.h5

Epoch 2/10

- 91s - loss: 2.2842 - acc: 0.3823 - val_loss: 2.1456 - val_acc: 0.4719

Epoch 00002: val_acc improved from 0.34502 to 0.47190, saving model to LSTM.h5

Epoch 3/10

- 91s - loss: 2.1144 - acc: 0.4649 - val_loss: 2.0799 - val_acc: 0.4702

Epoch 00003: val_acc did not improve from 0.47190

Epoch 4/10

- 91s - loss: 2.0793 - acc: 0.4668 - val_loss: 2.0760 - val_acc: 0.4252

Epoch 00004: val_acc did not improve from 0.47190

Epoch 5/10

- 91s - loss: 2.0536 - acc: 0.4717 - val_loss: 2.0572 - val_acc: 0.4613

Epoch 00005: val_acc did not improve from 0.47190

Epoch 6/10

- 91s - loss: 2.0729 - acc: 0.4757 - val_loss: 2.0762 - val_acc: 0.4636

Epoch 00006: val_acc did not improve from 0.47190

Epoch 7/10

- 91s - loss: 2.0500 - acc: 0.4733 - val_loss: 2.0107 - val_acc: 0.4769

Epoch 00007: val_acc improved from 0.47190 to 0.47691, saving model to LSTM.h5

Epoch 8/10

- 91s - loss: 2.0108 - acc: 0.4807 - val_loss: 1.9712 - val_acc: 0.4791

Epoch 00008: val_acc improved from 0.47691 to 0.47913, saving model to LSTM.h5

Epoch 9/10

- 90s - loss: 1.9944 - acc: 0.4793 - val_loss: 1.9784 - val_acc: 0.4780

Epoch 00009: val_acc did not improve from 0.47913

Epoch 10/10

- 91s - loss: 2.0023 - acc: 0.4821 - val_loss: 2.0570 - val_acc: 0.4791

